

Learning Scheduling Models from Event Data

removed for review

Abstract

Solving scheduling problems requires two main ingredients: a model that captures the essence of the scheduling problem and an algorithm that provides optimal solutions based on the model. Creating a suitable model that represents the problem was recognized to be a knowledge intensive task even when the underlying system is well-understood. In this work, we aim at automating the process of modeling scheduling problems. To this end, we introduce a novel three-step methodology for learning scheduling models from data. The first step mines timed Petri nets from event logs that contain executions of past schedules including information on activities, timestamps and resources. In the second step, we provide a mapping of timed Petri nets into *basic scheduling problems*, a family of models that represents a wide spectrum of scheduling problems. The third step shows a translation between basic scheduling problems and constraint programming formulations, which enables us to solve the problem to optimality. Our approach provides an end-to-end solution to model learning, going from data logs to model-based optimal schedules without human intervention. To demonstrate the usefulness of the methodology by conducting a series of experiments in which we learn scheduling models from two types of data: (1) event logs generated from job-shop scheduling benchmarks and (2) real-world event logs that come from an outpatient hospital. For the benchmark logs, the approach reconstructed the underlying job-shop scheduling models and provided optimal solutions. For the hospital data, our learning solution provided new daily schedules, thus exhibiting its practical relevance.

1 Introduction

Modern scheduling algorithms are able to successfully solve problems with thousands activities, complex temporal constraints, and scarce resources. The main prerequisite to solving scheduling problems is obtaining a mathematical model that fits a textual or a mathematical description of the underlying problem. For example, given a textual description of a deterministic Job-Shop Scheduling Problem (JSSP), the scheduling

expert uses the formulation to construct a constrained-based model of the JSSP, and to solve it using existing state-of-the-art algorithms (Baptiste, Pape, and Nuijten 2001, Chapter 6). Eliciting problem formulations often requires extensive domain knowledge and scheduling expertise (Lallouet et al. 2010).

In this work, we identify and address the challenge of learning scheduling models automatically from data without the need to formulate the problem first. For instance, instead of a textual description of the JSSP, i.e., ‘...a JSSP consists of n jobs that are waiting for m machines ...’ (Baptiste, Pape, and Nuijten 2001), we propose the use of historical event data that contains recordings related to the underlying system in order to *automatically* create a mathematical representation of the problem. Automating the process of model learning has the potential to minimize the costs related to modeling efforts and to turn scheduling problem formulation into an evidence-based procedure.

This paper is a first step in the direction of learning scheduling models. The data that we require for our approach is readily available in the form of event logs recorded by modern information systems. We assume that every event associated with the execution of previous schedules (e.g., start of operation and end of operation) is recorded along with its corresponding timestamp, label and resource. Our solution comprises of two phases, namely learning and transformation.

For the learning phase, we utilize existing approaches in the maturing field of *process mining* (van der Aalst 2011) to learn a formal problem representation in the form of a timed Petri net (TPN) (Van der Aalst 1996). **add how process mining mines out TPNs** For the transformation phase, we provide an algorithm that validates whether the learned TPN can be transformed into a basic scheduling problem (BSP). If the transformation is feasible, the algorithm maps the TPN into its corresponding BSP. Lastly, we show that the BSP is equivalent to a constraint programming (CP) formulation that can be used to solve the problem. **be a bit more precise about our approach**

To demonstrate the usefulness of our approach, we provide a two-part experimental evaluation. In the first part, we learn models of publicly available JSSP

benchmarks using synthetically generated data. In the second part of the evaluation, we learn scheduling models from real-world data coming from a large outpatient cancer hospital in the United States. Given the data, our algorithms generate the underlying appointment scheduling problem that we consequently solve using constraint programming.

2 Background

In this section we present the preliminaries to our approach. Firstly, we define timed Petri nets, a special case of stochastic Petri nets as defined in (Haas 2002, Chapter 2), having deterministic activity durations and routing. Secondly, we define our data model in the form of events logs and briefly outline existing approaches for learning TPN models from event logs. We conclude the section with a definition of a general family of scheduling problems, namely *basic scheduling problems* (BSP).

2.1 Timed Petri Nets

Petri nets are procedural models for analyzing discrete-event dynamic systems that exhibit parallelism, synchronization and resource consumption. Formally, a timed Petri net (TPN) is defined as follows:

Definition 1 (Timed Petri net (TPN); TPN System). *A timed petri net \mathcal{N} is a tuple $\mathcal{N} = \langle E, E', P, F, \tau \rangle$ with,*

- *E being a finite set of transitions with $E' \subseteq E$ being a (possibly empty) set of timed transitions,*
- *P being a finite set of places,*
- *$F \subseteq E \times P \cup P \times E$ being the flow relation of the Petri net, and,*
- *$\tau \in (E' \rightarrow \mathcal{T})$ being a function that maps deterministic durations to timed transitions;*

A TPN is fully characterized as a pair (\mathcal{N}, m_0) with \mathcal{N} being the net and $m_0 : P \rightarrow \mathbb{N}$ being an initial marking of the net, which is a function that maps each place to the number of tokens that it contains. We use Figure X to demonstrate a TPN and its dynamic behavior. In the figure, we observe two jobs waiting in the queue place.... **Insert simple (not too simple) TPN example here to explain the dynamics** We denote $\bullet p$ ($\bullet e$) the set of transitions (places) that precede a place p (a transition e), and by $p\bullet$ the set of transitions (places) that succeed a place p (a transition e). Furthermore, we define F_P to be the set of incoming and outgoing flows the corresponds to a set of places P , i.e., $F_P = \{(x, y) \in F \mid x \in P \vee y \in P\}$.

For example, in Figure X $\bullet p_3$... **complete here**

TPNs were previously applied to model scheduling problems (Van der Aalst 1996; Lee and DiCesare 1994). However, they provide an inefficient solution platform for solving scheduling problems, since TPNs require a global search over the entire state-space of the underlying problem (Lee and DiCesare 1994). Therefore, previous literature on scheduling with Petri nets mostly focuses on heuristic solutions (Lee and DiCesare 1994).

2.2 Mining Timed Petri nets from Event Logs

Process mining is a rapidly evolving research field that is centered around developing methodologies for learning models from data (van der Aalst 2011). The assumption is that the execution of processes is recorded into event logs, which can in turn be employed to learn models of the underlying system (e.g., in the form of a Petri net). The learning task is typically assumed to be unsupervised, since the learned model cannot be validated against ground truth.

To define process learning, we must first give an overview of our data model, namely the event log. Let \mathcal{E} be the universe of events with $e \in \mathcal{E}$ having attributes $e.j$ for job identifier, $e.s$ for start of operation timestamp, $e.c$ for completion of operation timestamp, $e.R$ for resource and $e.A$ for event label (e.g., operation 3 start, operation 5 complete, operation 7 start). An event log L is a subset of \mathcal{E} . Figure Y presents an excerpt from a real-world event log of an outpatient cancer hospital. Furthermore, we let $\psi(L, M) \in [0, 1]$ be a learning quality function that given an event log L and a TPN M evaluates the model with 0 (1) indicating low (high) quality model with respect to the log.

The task of a process learning function γ that maps an event log L onto a Petri net model M (in our case a TPN) such that $\psi(L, \gamma(L))$ is minimized (van der Aalst 2011). The measure ψ measures the distance between model and log indirectly, since we do not assume to have labeled pairs of logs and models. Many process learning algorithms were proposed in the past. See Chapter X in (van der Aalst 2011) for a survey.

In this work, we learn TPNs using the approach developed for scheduled processes (Senderovich et al. 2015). Specifically, the method in (Senderovich et al. 2015) learns the various components of the TPN, while guaranteeing maximal quality value, assuming that the event log was generated by a process that followed a pre-defined schedule. **Should we add more details here? no options between operations and loops are allowed**

3 Basic Scheduling Problem

In this part, we define a family of deterministic scheduling problems that we refer to as basic scheduling problems (BSPs). The BSP generalizes several well-known scheduling scheduling problems such as the job-shop scheduling problem (JSSP), and the resource-constrained project scheduling problem (RCPSP) (including its multi-mode variation). To define a BSP, we follow the scheduling problem formulation introduced by (Van der Aalst 1996).

Definition 2 (Basic Scheduling Problem (BSP)). *A basic scheduling problem is a tuple $\langle \mathcal{A}, \mathcal{R}, \Pi, c, d \rangle$ over a finite time domain \mathcal{T} with,*

- *\mathcal{A} being the set of activities to be scheduled,*
- *\mathcal{R} being the set of renewable resources,*



Figure 1: Our solution to learning scheduling models.

- $\Pi \subseteq \mathcal{A} \times \mathcal{A}$ being the precedence relation between pairs of activities,
- $c : \mathcal{R} \rightarrow \mathbb{N}^+$ being the function that maps resources to their capacities, and,
- $d : \mathcal{A} \times 2^{\mathcal{R}} \rightarrow \mathcal{T}$ being the duration partial function that maps pairs of activities and resource sets (that can execute these activities) to values in the time domain.

A schedule s , which is a solution to the BSP, is an allocation of resource sets to activities over time, i.e., $s \in \mathcal{A} \rightarrow (2^{\mathcal{R}} \times \mathcal{T})$. A feasible schedule respects resource and precedence constraints. Without loss of generality, one often considers an objective function $\phi(s) \in \mathbb{R}^{+0}$ that assigns a real-valued number to a given schedule. In optimal scheduling one aims at finding a schedule that minimizes $\phi(s)$.

A prominent example for a BSP is the *job shop scheduling problem* (JSSP), where only a single resource can perform each activity (hence $d(a), a \in \mathcal{A}$ is sufficient to represent durations), and resource capacities are equal to 1. The objective function in JSSP is to minimize the makespan.

Define the problem, present an overview of the approach

4 Solution Overview

In this section, we provide our end-to-end solution for learning scheduling models from data, as illustrated in Figure 1. Given an event log, we apply a process mining method from () to discover a timed Petri net, which represents the underlying system.

5 Learning Basic Scheduling Problems

In this section, we provide our approach to solving the schedule learning problem. We assume that a TPN is learned from event data using existing approaches, e.g., via the approach presented in (Senderovich et al. 2015). **we actually make a modification that enables that some of the activities can be performed by alternating sets of resources.**

5.1 Seize-Delay-Release Nets

The transformation of a learned TPN into a BSP, is based on the notion of Seize-Delay-Release nets with resources (SDRR nets).

To define SDRR nets, we first consider seize-delay-release constructs (SDCs), which are TPNs that consist of the following three components: (1) two transitions: one immediate transition that seizes a token and one timed transition that releases the token after a delay, (2) a place where the token is delayed, and (3) two

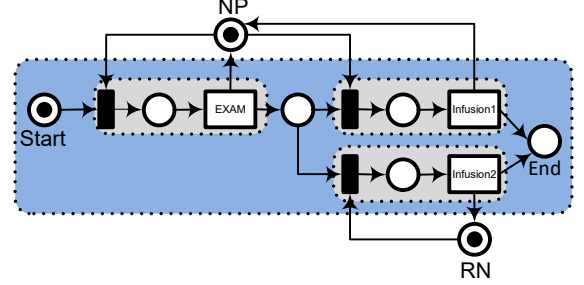


Figure 2: An SDRR Net of a Hospital Process.

flows that connect the two transitions to the delay place. Formally,

Definition 3 (Seize-Delay Construct (SDC)). *An SDC is a timed Petri net, $\mathcal{S} = \langle E, E', P, F, \tau \rangle$, such that*

- The set $E = \{e_{\text{seize}}, e_{\text{release}}\}$ contains two transitions (seize and delay),
- The set $E' = \{e_{\text{release}}\}$ is the timed delay transition,
- The set $P = \{p_{\text{delay}}\}$ is a single delay place, and,
- The flow is $F = \{(e_{\text{seize}}, p_{\text{delay}}), (p_{\text{delay}}, e_{\text{release}})\}$.

The gray parts in Figure 2 are the three SDC components of the TPN. The SDC that directly follows the start place in Figure 2 contains an immediate transition, which seizes the token in the start place. Then, the token is delayed for $\tau(e_{\text{release}})$, which corresponds to the duration of the ‘Exam’ transition. Lastly, ‘Exam’ releases the token into the subsequent place.

Given an SDC, \mathcal{S} , we denote $E_{\mathcal{S}}$ (and $E'_{\mathcal{S}}$), $P_{\mathcal{S}}$, $F_{\mathcal{S}}$ its sets of transitions (and timed transitions), places and flows, respectively. Furthermore, the set of places that precede (follow) the immediate (timed) transition of \mathcal{S} is denoted by $\bullet E_{\mathcal{S}}$ ($E_{\mathcal{S}} \bullet$), i.e.,

$$\begin{aligned} \bullet E_{\mathcal{S}} &= \{p \in P \mid \forall e \in E_{\mathcal{S}} \setminus E'_{\mathcal{S}} : p \in \bullet e\} \\ E_{\mathcal{S}} \bullet &= \{p \in P \mid \forall e' \in E'_{\mathcal{S}} : p \in e' \bullet\}. \end{aligned}$$

Detecting the SDCs of a given TPN is linear in the number of transitions ($\mathcal{O}(|E| + |P|)$), since it involves traversing over the immediate transitions of the TPN and verifying that the transition is followed by a single place, which is in turn followed by a single timed transition (see Definition 3).

A TPN that comprises a set of SDCs is referred to as a seize-delay-release net (SDR net). To define an SDR net, we let $S = \{\mathcal{S}_1, \dots, \mathcal{S}_m\}$ be a set of SDCs. The set S can be partitioned into k sets denoted by $C = \{\mathcal{C}_1, \dots, \mathcal{C}_k\}$, with each set $\mathcal{C}_j, j = 1, \dots, k$ containing SDCs that have

equal input and output places. It is easy to show that C always exists and that it is unique for a given TPN. To simplify notation, we denote E_{C_j} (E'_{C_j}) the transitions (timed transitions) of the SDCs that are elements of C_j . We are now ready to define seize-delay-release nets.

Definition 4 (Seize-Delay-Release Net). *A seize-delay-release net is a timed Petri net, $N_{sdr} = (E, E', P, F, \tau)$, which satisfies the following conditions:*

- The set $\bigcup_{j=1}^m E_{S_j}$ contains only transitions from the set of SDCs ($E' \subseteq E$ being the set of timed transitions),
- The set $P = \bigcup_{j=1}^m P_{S_j} \cup P_c$ contains both the delay places of N and a finite set of $k+1$ connector places $P_c = \{p_1, \dots, p_{k+1}\}$ with $p_1, p_{k+1} \in P_c$ being unique source and sink places, respectively, and,
- The flow $F = \bigcup_{j=1}^m F_{S_j} \cup F_c$ contains both the set of SDC flows and a set F_c such that

$$F_c = \{(p, e) \in P_c \times E \setminus E' \mid \quad (1)$$

$$p = p_j \wedge \exists C_j \in C (e \in E_{C_j} \setminus E'_{C_j}) \cup \{(e, p) \in E' \times P_c \mid \quad (2)$$

$$\exists C_j \in C (e \in E'_{C_j} \wedge p = p_{j+1})\}. \quad (3)$$

An example for an SDR net can be found in Figure 2. The blue area corresponds to an SDR net that consists of three SDCs partitioned into two sets, namely an SDC that involves ‘Exam’ and two SDCs that correspond to Infusion (transitions Infusion1 and Infusion2 are part of the same set in the partition of SDCs, C). Furthermore, the place of connectors P_c contains three places: start place, end place and a place that connects ‘Exam’ to the two ‘Infusion’ SDCs. SDR nets are important building blocks of the Seize-Delay-Release nets with resources (SDRR nets), which we map into basic scheduling problems. Algorithm 1 verifies that a TPN is an SDR net.

update explanation and proof

Below, we explain the algorithm line-by-line. Lines 2-5 verify that the TPN has unique input and output places. Line 6 returns the set of SDCs, S , that are present in the input TPN. Detecting SDCs is performed via the function *DetectSDC*(\cdot), which uses a simple traversal over all immediate transitions $e \in E \setminus E'$ and checking whether the direct followers of e are a single place and a timed transition (as in Definition 3). Lines 7-8 make sure that the set of transitions E includes only transitions from S .

Line 9 partitions the SDCs into sets of SDCs, such that each set $C \in C$ contains only SDCs that share input and output places.

Lines 10-15 traverse the sets in C and verify that each set $C_j \in C$ has a unique input place and a unique output place ($|\bullet C_j| = |C_j \bullet| = 1$). Furthermore, Lines 10-13 ensure that the input and output places of C_j are not equal to each other (no loops), and that there does not exist an additional set $C_i \neq C_j$ for which one of its input or output places are equal to the corresponding input and output places of C_j . Intuitively, this enforces

Algorithm 1: Verifies that TPN is SDR net.

Input: Timed Petri net $\mathcal{N} = \langle E, E', P, F, \tau \rangle$

Output: $\langle \{S_1, \dots, S_m\}, P_c \rangle$ - Set of SDCs and set of connector places P_c

```

1 begin
2    $P_{in} \leftarrow \{p \in P \mid \bullet p = \emptyset\}$ 
3    $P_{out} \leftarrow \{p \in P \mid p \bullet = \emptyset\}$ 
4   if  $|P_{in}| \neq 1 \vee |P_{out}| \neq 1$  then
5     return False
6    $S = \{S_1, \dots, S_m\} \leftarrow \text{DetectSDC}(\mathcal{N})$ 
7   if  $\bigcup_{j=1}^m E_{S_j} \neq E$  then
8     return False
9    $C \leftarrow \{C \subseteq S \mid \forall S_i, S_j \in C : \bullet E_{S_i} = \bullet E_{S_j} \wedge E_{S_i} \bullet = E_{S_j} \bullet\}$ 
10  foreach  $C_j \in C$  do
11     $\bullet C_j = \bigcup_{S \in C_j} \bullet E_S$ 
12     $C_j \bullet = \bigcup_{S \in C_j} E_S \bullet$ 
13    if  $(\bullet C_j = C_j \bullet \vee |\bullet C_j| \neq 1 \vee |C_j \bullet| \neq 1)$ 
14      then
15        return False
16  end
17  if  $\exists C_i \neq C_j \in C (C_i \bullet = C_j \bullet \vee \bullet C_i = \bullet C_j)$ 
18    then
19      return False
20   $P_c \leftarrow \bigcup_{C_j \in C} (\bullet C_j \cup C_j \bullet)$ 
21   $F_c \leftarrow \bigcup_{C_j \in C} \{(x, y) \in F \mid (x \in \bullet C_j \wedge y \in E_{C_j} \setminus E'_{C_j}) \vee (x \in E'_{C_j} \wedge y \in C_j \bullet)\}$ 
22  if  $P_c \neq P \setminus \bigcup_{j=1}^m P_{S_j} \vee F_c \neq F \setminus \bigcup_{j=1}^m F_{S_j}$ 
23    then
24      return False
25  return True
26 end
```

a sequential execution of the sets of SDCs in C . Lines 14-17 check whether the TPN places can be either SDC places or connectors (P_c) and that the only flows allowed in the TPN are within SDCs or between connector places, P_c , and SDC constructs.

If the algorithm reaches Line 18 without breaking, it concludes that the TPN is an SDR net and returns its set of SDCs S and the set of connector places P_c . Proposition 1 states the correctness and completeness of Algorithm 1. **polish proof**

Proposition 1. *Algorithm 1 is correct and complete, i.e., the algorithm returns non-empty sets, if and only if, the input TPN is an SDR net.*

Proof. If: Given an input TPN that is an SDR net, the algorithm returns non-empty sets.

The input TPN is an SDR net. Hence, by Definition 4, it has a source and a sink. Furthermore, $E = \bigcup_{j=1}^m E_{S_j}$ and therefore Lines 4 and 7 return True. In an SDR net, all flows are either in $\bigcup_{j=1}^m F_{S_j}$ (being part of an SDC) or in F_c , i.e., connecting places in $P_c = \{p_1, \dots, p_{k+1}\}$ to the set of SDCs in the partition set, C . Furthermore,

from the definition of F_c each partition set \mathcal{C}_j has exactly a single direct predecessor place $p_j \in P_c$ and a single direct successor place p_{j+1} . Moreover, no two sets in C can have the exact same predecessors and successors (otherwise, they would not be two different sets in C). Therefore, the condition in Line 13 holds for every $\mathcal{C}_j \in C$.

and only If: Given that the algorithm returns non-empty sets, the input net is an SDR net.

We prove the second direction by using the intermediate computations of Algorithm 1 to construct an SDR net $\mathcal{N} = (E, E', P, F, \tau)$ and showing that the net \mathcal{N} is equal to the input TPN. \square

Having defined SDR nets and an algorithm that verifies whether a given TPN is an SDR net, we define the seize-delay-release net with resources (SDRR net), which is a timed Petri net that consists of a set of SDR nets and a set of places and flows that correspond to resources. Let $N_{sdr} = \{\mathcal{N}_1, \dots, \mathcal{N}_n\}$ be a set of SDR nets with $\mathcal{N}_j = \langle E_j, E'_j, P_j, F_j, \tau_j \rangle, j = 1, \dots, n$ and let $S = \{\mathcal{S}_1, \dots, \mathcal{S}_m\}$ be a set of SD constructs that participate in N_{sdr} . We are now ready to define the SDRR nets.

Definition 5 (Seize-Delay-Release Net with Resources (SDRR net)). *An SDR net with resources (SDRR net) is a timed Petri net, $\mathcal{N}_{sdr} = (E, E', P, F, \tau)$, such that*

- The set $E = \bigcup_{j=1}^n E_j$ contains only transitions from the SDNs ($E' \subseteq E$ being the set of timed transitions),
- The set $P = \bigcup_{j=1}^m P_j \cup P_r$ contains both places from N and a finite set of resource places P_r , and,
- The flow set $F = \bigcup_{i=1}^m F_i \cup F_r$ contains all SDN flows and a set F_r such that:

$$F_r = \{(x, y) \in (P_r \times E \setminus E') \cup (E' \times P_r) \mid \forall (x, y) \in F_r : Q(x, y, S)\}$$

with,

$$\begin{aligned} Q(x, y, S) = & ((x \in P_r \wedge y \in E_{S_i} \setminus E'_{S_i} \Rightarrow \exists (e', x) \in F_r : e' \in E'_{S_i}) \\ & \wedge (x \in E'_{S_i} \wedge y \in P_r \Rightarrow \exists (y, e) \in F_r : e \in E_{S_i} \setminus E'_{S_i})). \end{aligned}$$

The set of resource flows F_r allows for resource tokens to be consumed only by immediate transitions and produced only by timed transitions. Furthermore, the property $Q(x, y)$ makes sure that if an immediate transition that consumes a resource token is part of an SDC, say \mathcal{S}_i , then there must also be a flow between the timed transition of \mathcal{S}_i back to the same resource place. Similarly, if a timed transition of an SDC produces a resource token, there must be a flow between the resource place and the immediate transition of the same SDC. One can easily verify that Figure 2 is an SDRR net with a single SDR net, three SDCs and two resource places (NP for nurse practitioner and RN for registered nurse).

Algorithm 2: Verifies that TPN is SDR net with resources (SDRR net).

Input: Timed Petri net $\mathcal{N} = \langle E, E', P, F, \tau \rangle$

Output: True or False

```

1 begin
2    $S = \{\mathcal{S}_1, \dots, \mathcal{S}_m\} \leftarrow \text{DetectSDC}(\mathcal{N})$ 
3    $P_r \leftarrow \{p \in P \mid \forall (x, y) \in F_{\{p\}}(Q(x, y, S))\}$ 
4    $\mathcal{N}' = (E, E', P \setminus P_r, F \setminus F_{P_r}, \tau)$ 
5    $\{\mathcal{N}_i\}_{i=1}^n \leftarrow \text{ConnectedComponents}(\mathcal{N}')$ 
6   if  $\exists i \in [n] : \text{VerifySDR}(\mathcal{N}_i) = \text{False}$  then
7     | return False
8   return True
9 end

```

Next, we introduce Algorithm 2, which take a TPN as its input, and returns whether the TPN is an SDRR net.

update this part for the new algorithm We shall briefly go over the algorithm and explain its parts. Lines 2-7 serve us to identify a set of places P_r that are candidates to be the resource places of the resulting SDRR. In an SDRR, a necessary condition for every $p \in P_r$ is that when p is removed from the net (along with the incoming and outgoing flows $F_{\{p\}}$) there will still exist a path between SDR net sources to their corresponding sinks. Removing a non-resource place, i.e., a place from one of the SDR nets of the SDRR net, will result in a source without a path to its sink. Therefore, removing places from the input TPN and verifying that every source has a path to a sink (via the function *AllSourceToSink* that receives a TPN and checks whether every source has a path to a sink), enables us to find candidates for P_r . Then, in Line 8, we remove all places in P_r from the input TPN and store the resulting net \mathcal{N}' . In Line 9, we store the connected components of \mathcal{N}' in a set of TPNs $\{\mathcal{N}_i\}_{i=1}^n$. If the input is an SDRR, this set will contain n SDR nets. Therefore, in Lines 12-18, we verify that the n connected components are SDR nets. Verifying whether \mathcal{N}_i is an SDR net is performed by running a Breadth-First Search (BFS) algorithm that verifies the conditions in Definition 4. The procedure *VerifySDR*(\cdot) returns \emptyset if \mathcal{N}_i is not an SDR (which will lead to an answer False in Algorithm 2 due to a violation of Definition 5); otherwise, it returns the set $P_{c,i}$, which is the connector set of the now verified SDR net \mathcal{N}_i , and S_i , which is the set of SDCs that \mathcal{N}_i comprises. Lastly, Lines 19-20 check whether the set of candidate resource places, P_r , satisfies $Q(x, y, S)$ from Definition 5. If one of the flows into $p \in P_r$ violates $Q(x, y, S)$, the algorithm returns False. Otherwise, the algorithm returns True.

polish proof

Proposition 2. *Algorithm 2 is correct and complete, i.e., the algorithm returns True, if and only if, the input to the algorithm is an SDRR net.*

Proof. If: Given an input TPN that is an SDRR net,

the algorithm returns True.

If \mathcal{N} is an SDRR net, then the only places that would not interrupt the flow between every source and sink of its SDR nets in N_{sdr} , are places in P_r . Therefore, the set P_r computed by the algorithm will contain only resource places of the SDRR net. After removing the resource places and their corresponding flows (Line 8), we are remained with n connected components, with each component being an SDR net (Line 9). The algorithm will verify the following two conditions: (1) the n components are SDR nets (Lines 10-18), and (2) the set of flows F_r respects $Q(x, y, S)$ (since the TPN is an SDRR net). Therefore it will return True.

and only If: Given that the algorithm returns True, the input net is an SDRR net.

We prove the second direction by using the intermediate computations of Algorithm 1 to construct a Petri net $\mathcal{N} = (E, E', P, F, \tau)$ and showing that \mathcal{N} corresponds to an SDRR net (Definition 5). We compose the SDRR net by connecting the TPNs $\mathcal{N}_i, i = 1, \dots, n$ (which are verified to be SDR nets by Lines 10-18) to the set of places P_r using the flow set F_r . Since the flows F_r respect property $Q(x, y)$, the resulting net is an SDRR net according to Definition 5. \square

5.2 Mapping SDRR nets into BSPs

In this part, we provide a mapping from Seize-Delay-Release timed Petri nets with resources (SDRR nets) into basic scheduling problems (BSPs). Specifically, given an SDRR net, $\mathcal{N} = (E, E', P, F, \tau)$, our we propose a mapping that creates the corresponding BSP tuple, $\langle \mathcal{A}, \mathcal{R}, \Pi, c, d \rangle$.

For conciseness, we refer to elements of the TPN (e.g., places, transitions, flows) when defining the BSP instead of labeling the TPN and then using these labels to define the BSP. **not sure that this needs to be said here**

For an SDRR net, we may reuse parts of Algorithms 1 and 2 to derive the following sets: (1) the set of SDCs $S = \{S_1, \dots, S_m\}$, (2) the set of resource places P_r , (3) the set of SDR nets that comprise the SDRR net $N_{sdr} = \{\mathcal{N}_i\}_{i=1}^n$ (4) the set of connector places $P_c = \{P_{c,i}\}_{i=1}^n$ for every SDR net that comprises the SDRR net, and (5) the partition of SDCs, $\{C_i\}_{i=1}^n$, which includes sets of SDCs with common input and output connectors in the i th SDR net.

Definition 6 (SDRR net to BSP Mapping). *Given an SDRR net $(\mathcal{N} = (E, E', P, F, \tau), m_0)$ and the sets $S, P_r, N_{sdr}, P_{c,i}, C_i$, the BSP is constructed as follows:*

- The resource set is given by the set of places, $\mathcal{R} = P_r$,
- The activity set corresponds to all connector places in P_c except the sinks, namely $\mathcal{A} = \{p \in P_c \mid p \bullet \neq \emptyset\}$,
- The precedence relation Π is constructed using the following:

$$\Pi = \{(a, b) \in \mathcal{A} \times \mathcal{A} \mid a \rightsquigarrow b\},$$

with \rightsquigarrow indicating that there exists a path from a to b ,

- resource capacities, c , are equal to the initial marking of resource places, i.e., $\forall r \in \mathcal{R} : c(r) = m_0(p_r)$, and finally,
- the duration (partial) function d is computed as follows:

$$d = \{(a, R, d) \in \mathcal{A} \times 2^{\mathcal{R}} \times \mathcal{T} \mid \\ \forall S \in S (\forall e \in E_S \setminus E'_S (a \in \bullet e \wedge R \subseteq (\bullet e \cap \mathcal{R})) \\ \wedge \forall e' \in E'_S (d = \tau(e')))\}$$

explain the mapping, prove its correctness?

5.3 Solving BSP with Constraint Programming

Kyle, please insert BSP to CP here. Chris will add a part on the intuition for scheduling people

6 Evaluation

7 Related Work

8 Conclusion

References

- [Baptiste, Pape, and Nuijten 2001] Baptiste, P.; Pape, C. L.; and Nuijten, W. 2001. *Constraint-Based Scheduling*. Norwell, MA, USA: Kluwer Academic Publishers.
- [Haas 2002] Haas, P. J. 2002. *Stochastic Petri Nets: Modelling, Stability, Simulation*. Springer.
- [Lallouet et al. 2010] Lallouet, A.; Lopez, M.; Martin, L.; and Vrain, C. 2010. On learning constraint problems. In *Tools with Artificial Intelligence (ICTAI), 2010 22nd IEEE International Conference on*, volume 1, 45–52. IEEE.
- [Lee and DiCesare 1994] Lee, D. Y., and DiCesare, F. 1994. Scheduling flexible manufacturing systems using petri nets and heuristic search. *IEEE Transactions on robotics and automation* 10(2):123–132.
- [Senderovich et al. 2015] Senderovich, A.; Rogge-Solti, A.; Gal, A.; Mendling, J.; Mandelbaum, A.; Kadish, S.; and Bunnell, C. A. 2015. Data-driven performance analysis of scheduled processes. In Motahari-Nezhad, H. R.; Recker, J.; and Weidlich, M., eds., *Business Process Management - 13th International Conference, BPM 2015, Innsbruck, Austria, August 31 - September 3, 2015, Proceedings*, volume 9253 of *Lecture Notes in Computer Science*, 35–52. Springer.
- [Van der Aalst 1996] Van der Aalst, W. 1996. Petri net based scheduling. *Operations-Research-Spektrum* 18(4):219–229.
- [van der Aalst 2011] van der Aalst, W. M. P. 2011. *Process Mining - Discovery, Conformance and Enhancement of Business Processes*. Springer.